

MOHAMMED ANSARI

San Jose, California

☎ (669) 288-3655

✉ mohammed.ansari@sjsu.edu

🌐 [linkedin.com/in/imohammedansari](https://www.linkedin.com/in/imohammedansari)

🌐 mohammedansari.dev

🐙 github.com/ByteMaverick

Education

San Jose State University

Aug. 2023 – May 2027

Bachelor of Science in Data Science — GPA: 3.94(Spring 25) | Cumulative: 3.56

San Jose, California

- **Coursework:** Data Structures and Algorithms, Machine Learning, Advanced Python, Processing Big Data, Intro to Artificial Intelligence, Object Oriented Design

Technical Skills

Programming Languages: Python, Java, Bash Scripting

Machine Learning & LLMs: Scikit-Learn, TensorFlow, Hugging Face, OpenAI and Vertex AI APIs

XAI & Data Analysis: SHAP, LIME, Evidently AI, Pandas, NumPy, Matplotlib, Seaborn

Databases & Monitoring: MySQL, MongoDB, Vector Databases, Grafana, Prometheus

Cloud & MLOps: AWS, Google Cloud, Docker, Apache Airflow, MLflow, Git, GitHub Actions(CI/CD)

Concepts & Frameworks: RAG, Vector Search, Prompt Engineering, A/B Testing, FastAPI, LangChain, LangGraph

Technical Experience

ClauseBit – AI-Powered Privacy Analyzer 🔗

May 2025 – July 2025

Creator

San Jose, California

- Designed an **AI-powered browser extension + web app** to analyze and summarize Privacy Policies and Terms of Service using **user-defined privacy preferences**; led full-stack development and Figma prototyping.
- Developed an **Adaptive RAG pipeline** using **LLM APIs + Vector Search** to identify and explain **risky legal clauses** (e.g., data sharing, dark patterns), enabling **real-time clause-level QA** and summaries.
- Built a self-updating **agentic knowledge base** with on-demand web scraping agents, cutting manual maintenance by **60%** and enriching sparse documents.
- Reduced **hallucinations by 35%** using clause relevance filtering, vector confidence scoring, and fallback rejection.
- Deployed full-stack system on **GCP (Cloud Run, Functions, Firestore)** with **Clerk** authentication, optimizing end-to-end latency to **sub-20s** using chunking, LRU caching, and **async-based concurrency**.

DockFlow: Real-Time Bike Dock Availability Predictor 🔗

Aug 2025 – Present

Creator

San Jose, California

- Built a real-time ML platform to forecast Bay Wheels docking station availability in San Jose within 15 minutes by integrating **GBFS live feeds**, **weather data**, and **event schedules** into a unified, time-partitioned dataset.
- Orchestrated a robust **Apache Airflow** ETL pipeline on **GCS** and **BigQuery** to automate multi-source ingestion, cleaning, and feature engineering, ensuring reliable and reproducible workflows.
- Implemented continuous **model drift monitoring** with **Evidently AI**, applying **Kolmogorov–Smirnov** and **Population Stability Index** tests to detect performance decay and trigger adaptive retraining.
- Deployed a containerized **FastAPI** inference service to **GCP Cloud Run** with CI/CD via **GitHub Actions** and comprehensive monitoring in **Prometheus & Grafana** for model, pipeline, and system health.

TripSafe: Traffic Accident Risk Prediction Tool 🔗

April 2025 – May 2025

Team Lead

San Jose, California

- Performed **EDA** and led end-to-end development of a **multi-class accident risk classifier** using 7M+ U.S. accident records, processed via **Bash scripting** for large-scale data cleaning and preparation.
- Defined interpretable risk labels (Low, Medium, High) via **K-Means** clustering; applied **PCA** to reduce dimensionality.
- Balanced class distribution with **SMOTE**; tuned **Logistic Regression**, **SVM**, **XGBoost** via **GridSearchCV**.
- Achieved **95% test accuracy** via **cross-validation** with low log loss; validated generalization using **classification reports** and **confusion matrices** for deployment in driver safety and route optimization systems.

Leadership/Extracurricular

Machine Learning Club @ SJSU

Feb 2025 – Present

Lead Outreach and Recruitment Officer

San Jose, California

- Led onboarding and outreach campaigns that increased member retention by **25%** and drove consistent event turnout.
- Monitored and analyzed weekly **engagement metrics** to optimize recruitment strategies and boost participation.

Certifications and Awards

• AWS Certified AI Practitioner

Jun 2025

• Machine Learning A-Z (Udemy)

Jul 2025

• Dean's List, San Jose State University

Fall 2023 – Spring 2024