

MOHAMMED ANSARI

San Jose, California

📞 (669) 288-3655

✉️ mohammed.ansari@sjtu.edu

🔗 linkedin.com/in/imohammedansari

GitHub github.com/ByteMaverick

🌐 mohammedansari.dev

Education

San Jose State University

Bachelor of Science in Data Science

Aug. 2023 – May 2027

San Jose, California

- Coursework: Advanced Machine Learning, Advanced Python, Big Data, Object Oriented Design, DSA, Statistics

Technical Skills

Languages: Python, Java, SQL, Bash

Machine Learning / AI: PyTorch, TensorFlow, Scikit-Learn, Hugging Face, LangChain, OpenAI API

Data Science & Analytics: Pandas, NumPy, Scipy, Seaborn, Plotly, Panels

GenAI & Retrieval: RAG, Vector Search (FAISS, Pinecone), Transformers, NLP, Prompt Engineering

Cloud & MLOps: AWS, GCP, Docker, Airflow, MLflow, Evidently AI, FastAPI, CI/CD (GitHub Actions)

Experience

Spartan Knowledge-Driven Intelligence & Learning Lab @ SJSU

Sep 2025 – Present

Undergraduate Researcher

San Jose, California

- Collaborated with faculty to develop a **baseline architecture** and **reasoning pipeline** that integrated LLMs with Knowledge Graph search to enhance factual transparency and interpretability in enterprise-scale AI systems.
- Investigating **coverage gaps** in Knowledge Graphs by integrating link prediction with LLM-based validation to identify and verify missing relations, ensuring more accurate and trustworthy knowledge-driven reasoning.

Artificial Intelligence & Machine Learning Club @ SJSU

Feb 2025 – Present

Lead Outreach Officer

San Jose, California

- Led **5-member** team hosting generative-AI and responsible-ML workshops for 200+ students.
- Collaborate with faculty and industry speakers to design hands-on sessions to foster campus-wide engagement in AI/ML.

Projects

DockFlow – Real-Time Bike Dock Availability Predictor

Aug 2025 – Present

Creator

San Jose, California

- Developed a real-time **ML system** predicting Bay Wheels dock availability 15 min ahead using live data & weather.
- Orchestrated a scalable **Airflow ETL pipeline** on **GCS** for automated preprocessing and time-series model training.
- Deploying a containerized **FastAPI** service to **GCP Cloud Run** with CI/CD (**GitHub Actions**) and integrating **Evidently AI** for continuous model-drift monitoring and automated retraining.

ClauseBit – AI-Powered Privacy Analyzer

May 2025 – July 2025

Creator

San Jose, California

- Built an **AI-driven browser extension** that automated privacy-policy analysis, reducing manual review time by **40%** compared to human baseline (avg. 5 → 3 min per policy).
- Designed an **adaptive RAG pipeline** using LLM APIs, vector search, and clause-relevance filtering to cut hallucinations by **40%** relative to base LLM outputs.
- Developed a self-updating **knowledge base** with web-scraping agents and Pydantic models, and deployed the full-stack system on **GCP Cloud Run & Firestore** with async caching for sub-25s latency.

TripSafe – Traffic Accident Risk Prediction Tool

Apr 2025 – May 2025

Team Lead — Course Project – A Grade

San Jose, California

- Led development of a **multi-class accident risk model** on **7M+ U.S. records** to identify high-risk routes and improve traffic-safety decision making.
- Created interpretable risk labels (Low, Medium, High) through **K-Means** clustering and applied **PCA** to reduce dimensionality for efficient, high-performing multi-class classification.
- Tuned **Random Forest**, **SVM**, **XGBoost** with GridSearchCV and achieving **95%** accuracy under cross-validation.

Certifications and Awards

AWS Certified AI Practitioner

Jun 2025

Dean's List, San Jose State University

Fall 2023 – Spring 2024